

# Storage

## S3 (Simple Storage Service)

S3 Bucket is an **object-based** storage, used to manage data as objects

**S3 Object** is having:-

Value - data bytes of object (photos, videos, documents, etc.)

Key - full path of the object in bucket e.g. `/movies/comedy/abc.avi`

Version ID - version object, if versioning is enabled

Metadata - additional information

S3 Bucket holds objects. S3 console shows virtual folders based on key.

S3 is a universal namespace so bucket names must be globally unique (think like having a domain name)

```
https://<bucket-name>.s3.<aws-region>.amazonaws.com  
or  
https://s3.<aws-region>.amazonaws.com/<bucket-name>
```

Unlimited Storage, Unlimited Objects from **0 Bytes** to **5 Terabytes** in size. You should use **multi-part upload** for Object size > **100MB**

All new buckets are **private** when created by default. You should enable **public access** explicitly.

Access control can be configured using **Access Control List (ACL)** (deprecated) and **S3 Bucket Policies** (recommended)

**S3 Bucket Policies** are **JSON** based policy for complex access rules at user, account, folder, and object level

Enable **S3 Versioning** and **MFA delete** features to protect against accidental delete of S3 Object.

Use **Object Lock** to store object using write-once-read-many (WORM) model to prevent objects from being deleted or overwritten for a fixed amount of time (**Retention period**) or indefinitely (**Legal hold**). Each version of object can have different retention-period.

You can **host static websites** on S3 bucket consisting of HTML, CSS, **client-side JavaScript**, and images. You need to enable Static website hosting and Public access for S3 to avoid 403 forbidden error. Also you need to add **CORS Policy** to allow cross-origin request.

```
https://<bucket-name>.s3-website[.-]<aws-region>.amazonaws.com
```

Generate a **pre-signed URL** from CLI or SDK (can't from the web) to provide temporary access to an S3 object to either upload or download object data. You specify expiry (say 5 sec) while generating url:-

```
aws s3 presign s3://mybucket/myobject --expires-in 300
```

**S3 Select** or **Glacier Select** can be used to query subset of data from S3 Objects using SQL query. S3 Objects can be CSV, JSON, or Apache Parquet. GZIP & BZIP2 compression is supported with CSV or JSON format with server-side encryption.

using **Range** HTTP Header in a GET Request to download the specific range of bytes of S3 object, known as **Byte Range Fetch**

You can create **S3 event notification** to push events e.g. `s3:ObjectCreated:*` to SNS topic, SQS queue or execute a Lambda function. It is possible that you receive single notification for two writes to a non-versioned object at the same time. Enable versioning to ensure you get all notifications.

Enable **S3 Cross-Region Replication** for asynchronous replication of object across buckets in another region. You must have **versioning** enabled on both **source** and **destination** side. Only **new S3 Objects** are replicated after you enable them.

Enable **Server access logging** for logging object-level fields object-size, total time, turn around time, and HTTP referrer. Not available with CloudTrail.

Use **VPC S3 gateway endpoint** to access S3 bucket within AWS VPC to reduce the overall data transfer cost.

Enable **S3 Transfer Acceleration** for faster transfer and high throughput to S3 bucket (mainly uploads), Create **CloudFront** distribution with OAI pointing to S3 for faster-cached content delivery (mainly reads)

Restrict the access of S3 bucket through CloudFront only using **Origin Access Identity (OAI)**. Make sure user can't use a direct URL to the S3 bucket to access the file.

## S3 Storage Class Types

**Standard:** Costly choice for very high availability, high durability and fast retrieval

**Intelligent Tiering:** Uses ML to analyze your Object's usage and move to the appropriate cost-effective storage class automatically

**Standard-IA:** Cost-effective for infrequent access files which cannot be recreated

**One-Zone IA:** Cost-effective for infrequent access files which can be recreated

**Glacier:** Cheaper choice to Archive Data. You must purchase **Provisioned capacity**, when you require guaranteed **Expedite retrievals**.

**Glacier Deep Archive:** Cheapest choice for Long-term storage of large amount of data for compliance

S3 Storage Class	Durability	Availability	AZ	Min. Storage	Retrieval Time	Retrieval fee
S3 Standard (General Purpose)	11 9's	99.99%	≥3	N/A	milliseconds	N/A
S3 Intelligent Tiering	11 9's	99.9%	≥3	30 days	milliseconds	N/A
S3 Standard-IA (Infrequent Access)	11 9's	99.9%	≥3	30 days	milliseconds	per GB
S3 One Zone-IA (Infrequent Access)	11 9's	<b>99.5%</b>	<b>1</b>	30 days	milliseconds	per GB

S3 Storage Class	Durability	Availability	AZ	Min. Storage	Retrieval Time	Retrieval fee
S3 Glacier	11 9's	99.99%	≥3	90 days	<b>Expedite (1-5 mins)</b> Standard (3-5 hrs) Bulk (5-12 hrs)	per GB
S3 Glacier Deep Archive	11 9's	99.99%	≥3	180 days	<b>Standard (12 hrs)</b> Bulk (48 hrs)	per GB

You can upload files in the same bucket with different **Storage Classes** like *S3 standard*, *Standard-IA*, *One Zone-IA*, *Glacier* etc.

You can setup **S3 Lifecycle Rules** to transition current (or previous version) objects to cheaper storage classes or delete (expire if versioned) objects after certain days e.g. transition from S3 Standard to S3 Standard-IA or One Zone-IA can only be done after 30 days.

transition from S3 Standard to S3 Intelligent Tiering, Glacier, or Glacier Deep Archive can be done immediately.

You can also setup lifecycle rule to **abort multipart upload**, if it doesn't complete within certain days, which auto delete the parts from S3 buckets associated with multipart upload.

## Encryption

**Encryption in transit** between client and S3 is achieved via SSL/TLS

You can add default encryption at bucket level and also override encryption at file level.

### Encryption at rest - Server Side Encryption (SSE)

**SSE-S3** AWS S3 managed keys, use AES-256 algorithm. Must set header: `"x-amz-server-side-encryption":"AES-256"`

**SSE-KMS** Envelope Encryption using AWS KMS managed keys. Must set header: `"x-amz-server-side-encryption":"aws:kms"`

**SSE-C** Customer provides and manage keys. HTTPS is mandatory.

**Encryption at rest - Client Side Encryption** client encrypts and decrypts the data before sending and after receiving data from S3.

To meet **PCI-DSS or HIPAA** compliance, encrypt S3 using SSE-C and Client Side Encryption

## Data Consistency

S3 provides **strong read-after-write consistency for PUTs and DELETes** of objects. PUTs applies to both writes to new objects as well as overwrite existing objects.

**Updates to a single key are atomic.** For example, if you PUT to an existing key from one thread and perform a GET on the same key from a second thread concurrently, you will get either the old data or the new data, but never partial or corrupt data.

## AWS Athena

You can use **AWS Athena** (Serverless Query Engine) to perform analytics directly against S3 objects using SQL query and save the analysis report in another S3 bucket.

**Use Case:** one-time SQL query on S3 objects, S3 access log analysis, serverless queries on S3, IoT data analytics in S3, etc.

## Instance Store

Instance Store is temporary **block-based** storage physically attached to an EC2 instance  
Can be attached to an EC2 instance only when the instance is launched and cannot be dynamically resized

Also known as **Ephemeral Storage**

Deliver very low-latency and high random I/O performance

Data persists on instance reboot, data **doesn't persist on stop or termination**

## EBS (Elastic Block Store)

EBS is **block-based** storage, referred as EBS Volume

**EBS Volume** think like a USB stick

Can be attached to only one EC2 instance at a time. Can be detached & attached to another EC2 instance in that same AZ only

Can attach multiple EBS volumes to single EC2 instance. Data persist after detaching from EC2

**EBS Snapshot** is a backup of EBS Volume at a point in time. You can not copy EBS volume across AZ but you can create EBS Volume from Snapshot across AZ. EBS Snapshot can copy across AWS Regions.

Facts about EBS Volume **encryption**:-

All data at rest inside the volume is encrypted

All data in flight between the volume and EC2 instance is encrypted

All snapshots of encrypted volumes are automatically encrypted

All volumes created from encrypted snapshots are automatically encrypted

Volumes created from unencrypted snapshots can be encrypted at the time of creation

EBS supports **dynamic changes in live production** volume e.g. volume type, volume size, and IOPS capacity without service interruption

There are two types of EBS volumes:-

**SSD** for small/random IO operations, High IOPS means number of read and write operations per second, Only SSD EBS Volumes can be used as **boot volumes** for EC2

**HDD** for large/sequential IO operations, High Throughput means number of bytes read and write per second

EBS Volumes with two types of RAID configuration:-

**RAID 0** (increase performance) two 500GB EBS Volumes with 4000 IOPS - creates 1000GB RAID0 Array with 8000 IOPS and 1000Mbps throughput

**RAID 1** (increase fault tolerance) two 500GB EBS Volumes with 4000 IOPS - creates 500GB RAID1 Array with 4000 IOPS and 500Mbps throughput

EBS Volume Types	Description	Usage
General Purpose SSD (gp2/gp3)	Max 16000 IOPS	boot volumes, dev environment, virtual desktop
Provisioned IOPS SSD (io1/io2)	16000 - 64000 IOPS, EBS Multi-Attach	critical business application, large SQL and NoSQL database workloads
Throughput Optimized HDD (st1)	Low-cost, frequently accessed, throughput intensive	Big Data, Data warehouses, log processing

EBS Volume Types	Description	Usage
Cold HDD (sc1)	Lowest-cost, infrequently accessed	Large data with lowest cost

## EFS (Elastic File System)

EFS is a **POSIX-compliant file-based** storage

EFS supports **file systems semantics** - strong read-after-write consistency and file locking

**highly scalable** - can automatically scale from gigabytes to petabytes of data without needing to provision storage. With **burst mode**, the throughput increase, as file system grows in size.

**highly available** - stores data redundantly across multiple Availability Zones

Network File System (NFS) that can be mounted on and **accessed concurrently by thousands of EC2** in multiple AZs without sacrificing performance.

EFS file systems can be accessed by Amazon EC2 Linux instances, Amazon ECS, Amazon EKS, AWS Fargate, and AWS Lambda functions via a file system interface such as NFS protocol.

### Performance Mode:

**General Purpose** for most file system for low-latency file operations, good for content-management, web-serving etc.

**Max I/O** is optimized to use with 10s, 100s, and 1000s of EC2 instances with high aggregated throughput and IOPS, **slightly higher latency** for file operations, good for big data analytics, media processing workflow

**Use case:** Share files, images, software updates, or computing across all EC2 instances in ECS, EKS cluster

## FSx for Windows

Windows-based file system supports **SMB** protocol & Windows **NTFS**

supports **Microsoft Active Directory (AD) integration**, ACLs, user quotas

## FSx for Lustre

Lustre = Linux + Cluster is a **POSIX-compliant parallel linux file system**, which stores data across multiple network file servers

High-performance file system for **fast processing of workload** with consistent **sub-millisecond latencies**, up to hundreds of gigabytes per second of throughput, and up to **millions of IOPS**.

Use it for Machine learning, High-performance computing (HPC), video processing, financial modeling, genome sequencing, and electronic design automation (EDA).

You can use **FSx for Lustre as hot storage** for your highly accessed files, and **Amazon S3 as cold storage** for rarely accessed files.

**Seamless integration with Amazon S3** - connect your S3 data sets to your FSx for Lustre file system, run your analyses, write results back to S3, and delete your file system

FSx for Lustre provides two deployment options:-

**Scratch file systems** - for temporary storage and short-term processing

**Persistent file systems** - for high available & persist storage and long-term processing

